

Lab 5 - Python

Table of contents

1	Data Analysis	1
1.1	Table Presentation	1
1.2	Graphs - Final output	3
2	Linear Regression Analysis	5
2.1	LM 1	5
2.2	LM 2	6
2.3	Double Lasso	6
2.4	Results	7
3	Non Linear Methods	8
3.1	Lasso	8
3.2	Regression Trees	8
3.3	Boosting Trees	9
3.4	Regresssion Forest	9
3.5	Results	10
3.5.1	Plot	10

1 Data Analysis

1.1 Table Presentation

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
```

```
warnings.filterwarnings("ignore")

data = pd.read_csv(
    "https://raw.githubusercontent.com/alexanderquispe/CausalAI-Course/main/data/processed_e"
)
```

```
data1 = (
    data.filter(
        [
            "y",
            "w",
            "gender_female",
            "gender_male",
            "gender_transgender",
            "age",
            "imd_decile",
        ]
    )
    .melt(
        id_vars=["y", "w", "age", "imd_decile"],
        value_vars=["gender_female", "gender_male", "gender_transgender"],
        var_name="gender",
        value_name="value",
    )
    .query("value > 0")
    .assign(
        gender=lambda df: df["gender"]
        .str.split("_")
        .str[-1]
        .map({"female": "Female", "male": "Male", "transgender": "Transgender"})
    )
    .drop(columns=["value"])
)

summary = (
    data1.groupby(["w", "gender"])
    .agg(
        n=("y", "size"),
        mean_y=("y", "mean"),
        sd_y=("y", "std"),
        mean_age=("age", "mean"),
        sd_age=("age", "std"),
    )
)
```

```

)
.reset_index()
.assign(
    w=lambda df: df["w"].map({0: "Control", 1: "Treatment"}),
    gender=lambda df: df["gender"].map(
        {"Male": "Male", "Female": "Female", "Transgender": "Transgender"}
    ),
)
.rename(
    columns={
        "w": "Group",
        "gender": "Gender",
        "mean_y": "Mean - Y",
        "sd_y": "sd - Y",
        "mean_age": "Mean - Age",
        "sd_age": "sd - Age",
    }
)
)

print(summary)

```

	Group	Gender	n	Mean - Y	sd - Y	Mean - Age	sd - Age
0	Control	Female	475	0.208421	0.406608	22.741053	3.591231
1	Control	Male	342	0.216374	0.412376	23.491228	3.547530
2	Control	Transgender	1	0.000000	NaN	17.000000	NaN
3	Treatment	Female	541	0.534196	0.499291	22.920518	3.500154
4	Treatment	Male	377	0.389920	0.488380	23.506631	3.575267
5	Treatment	Transgender	3	1.000000	0.000000	22.333333	3.214550


1.2 Graphs - Final output

```


sns.set_theme(style="whitegrid")
plt.figure(figsize=(10, 6))
sns.countplot(data=data1, x='w', hue='gender', palette="Set2")
plt.xlabel('')
plt.ylabel('n')
plt.title('Gender Distribution')
plt.legend(title='Gender', loc='upper left')

```

```
plt.xticks([0, 1], ['Control', 'Treatment'])
plt.show()
```



```
plt.figure(figsize=(10, 6))
sns.kdeplot(data=data1, x='age', hue='w', fill=True, palette=['#0798a8', '#a8073a'], alpha=0.5)
plt.xlabel('Age')
plt.ylabel('Density')
plt.title('Density')
plt.legend(title='', loc='lower right', labels=['Control', 'Treatment'])
plt.show()
```



2 Linear Regression Analysis

2.1 LM 1

```
import statsmodels.api as sm
import statsmodels.formula.api as smf
lm1 = smf.ols('y ~ w', data=data).fit()
lm1_summary = lm1.summary2().tables[1]
lm1_summary['model'] = "Simple"
lm1_summary
```

	Coef.	Std.Err.	t	P> t	[0.025	0.975]	model
Intercept	0.211491	0.016053	13.174492	7.533982e-38	0.180006	0.242977	Simple
w	0.265164	0.022059	12.020880	4.957428e-32	0.221900	0.308429	Simple

2.2 LM 2

```
base_formula = 'y ~ w + age + imd_decile'
lm2 = smf.ols(base_formula, data=data1).fit()
lm2_summary = lm2.summary2().tables[1]
lm2_summary['model'] = "Multiple"
lm2_summary
```

	Coef.	Std.Err.	t	P> t	[0.025	0.975]	model
Intercept	-0.147526	0.077502	-1.903520	5.713801e-02	-0.299533	0.004481	Multiple
w	0.263377	0.021907	12.022335	4.892918e-32	0.220409	0.306344	Multiple
age	0.015804	0.003069	5.149001	2.916669e-07	0.009784	0.021824	Multiple
imd_decile	-0.001501	0.007416	-0.202367	8.396539e-01	-0.016046	0.013045	Multiple

2.3 Double Lasso

```
from sklearn.linear_model import LassoCV, Lasso
import numpy as np
cnt = [
    "w",
    "gender_female",
    "gender_male",
    "age",
    "imd_decile"
]

X = data[cnt]
y = data['y']

lasso_cv = LassoCV(alphas = None, cv = 4, max_iter=10).fit(X, y)
model_glm = Lasso(alpha = lasso_cv.alpha_).fit(X, y)
coefs = model_glm.coef_

relevant = np.nonzero(coefs)[0]


xvar= np.array(cnt)[relevant]
base_model = 'y ~ ' + " + ".join(xvar)
lm3 = smf.ols(base_model, data=data).fit()
lm3_summary = lm3.summary2().tables[1]
```

```
lm3_summary['model'] = "Lasso"
lm3_summary
```

	Coef.	Std.Err.	t	P> t	[0.025	0.975]	model
Intercept	-0.133280	0.077287	-1.724486	8.479844e-02	-0.284866	0.018305	Lasso
w	0.262487	0.021823	12.028213	4.589079e-32	0.219685	0.305288	Lasso
gender_male	-0.085201	0.022225	-3.833634	1.307738e-04	-0.128791	-0.041611	Lasso
age	0.016887	0.003070	5.500006	4.363906e-08	0.010865	0.022909	Lasso
imd_decile	-0.002528	0.007392	-0.341942	7.324357e-01	-0.017026	0.011970	Lasso

2.4 Results

```
combined_results = pd.concat([lm1_summary, lm2_summary, lm3_summary])
filtered_results = combined_results.query('index == "w"')
filtered_results.columns = ["coef", "std", "t", "p", "u", "l", "model"]
for model in filtered_results.model:
    ref = filtered_results.query("model == @model")
    x = [model, model]
    y = [ref["u"], ref["l"]]
    plt.scatter("model", "coef", data=ref)
    # plt.scatter('')
    plt.plot(x, y)
```



3 Non Linear Methods

```
from sklearn.model_selection import train_test_split

train, test = train_test_split(
    data, stratify=data['y'], test_size=0.3, random_state=23
)
y_train = train['y']
d_train = train['w']
x_train = train.drop(columns = ['y', 'w'])

y_test = test['y']
d_test = test['w']
x_test = test.drop(columns = ['y', 'w'])
```

```
def lm_yd(y, d, md: str):
    df = pd.DataFrame({'y': y, 'd': d})
    model = smf.ols("y ~ d", data = df).fit()
    summary = model.summary2().tables[1]
    summary['model'] = md
    summary.columns = ['estimate', 'std.error', 't', 'p', 'u', 'l', 'model']
    return summary
```

3.1 Lasso

```
def residual_lasso(X_train, y_train, X_test, y_test, cv_n = 10):
    model = LassoCV(cv = cv_n).fit(X_test, y_test)
    residual = y_test - model.predict(X_test)
    return residual

y_r1 = residual_lasso(x_train, y_train, x_test, y_test)
d_r1 = residual_lasso(x_train, d_train, x_test, d_test)
l1 = lm_yd(y_r1, d_r1, "Lasso")
```

3.2 Regression Trees


```

from sklearn.tree import DecisionTreeRegressor

def residual_dtree(X_train, y_train, X_test, y_test):
    model = DecisionTreeRegressor().fit(X_train, y_train)
    residual = y_test - model.predict(X_test)
    return residual

y_r2 = residual_dtree(x_train, y_train, x_test, y_test)
d_r2 = residual_dtree(x_train, d_train, x_test, d_test)

l2 = lm_yd(y_r2, d_r2, "Reg. Trees")

```

3.3 Boosting Trees

```

from sklearn.ensemble import GradientBoostingRegressor

def gbm_residual(xtrain, xtest, ytrain, ytest):
    model = GradientBoostingRegressor().fit(xtrain, ytrain)
    residual = ytest - model.predict(xtest)
    return residual

y_r3 = gbm_residual(x_train, x_test, y_train, y_test)
d_r3 = gbm_residual(x_train, x_test, d_train, d_test)

l3 = lm_yd(y_r3, d_r3, "Bost. Trees")

```

3.4 Regression Forest

```

from sklearn.ensemble import RandomForestRegressor

def rand_forest_res(x_train, y_train, x_test, y_test) :
    model = RandomForestRegressor().fit(x_train, y_train)
    residual = y_test - model.predict(x_test)
    return residual

y_r4 = rand_forest_res(x_train, y_train, x_test, y_test)
d_r4 = rand_forest_res(x_train, d_train, x_test, d_test)

```

```
l4 = lm_yd(y_r4, d_r4, "Reg. Forest")
```

3.5 Results

```
from tabulate import tabulate

results = pd.concat([l1, l2, l3, l4])
results = results.query("index == 'd'").sort_values("estimate")

print(tabulate(results.round(2), headers="keys", tablefmt="pretty", showindex=False))
```

estimate	std.error	t	p	u	l	model
0.17	0.04	4.05	0.0	0.09	0.25	Reg. Trees
0.19	0.04	4.97	0.0	0.12	0.27	Reg. Forest
0.21	0.04	5.4	0.0	0.14	0.29	Bost. Trees
0.25	0.04	6.24	0.0	0.17	0.33	Lasso

3.5.1 Plot

```
for md in results["model"]:
    ref = results.query("model == @md")
    plt.scatter("model", "estimate", data=ref)
    ci = ref[["u", "l"]].values[0]
    label = [md, md]
    plt.plot(label, ci)
```

