

Modern High Dimensional Linear Regression

Alexander Quispe
The World Bank, PUCP
alexander.quispe@pucp.edu.pe

April 3, 2024

Analysis of Variance (ANOVA)

POPULATION	SAMPLE
$Y = \beta'X + \varepsilon, \quad E\varepsilon X = 0$	$Y_i = \hat{\beta}'X_i + \hat{\varepsilon}_i$
$EY^2 = E(\beta'X)^2 + E\varepsilon^2$	$\mathbb{E}_n Y_i^2 = \mathbb{E}_n(\hat{\beta}'X_i)^2 + \mathbb{E}_n \hat{\varepsilon}_i^2$
$MSE_{pop} = E\varepsilon^2$	$MSE_{sample} = \mathbb{E}_n \hat{\varepsilon}_i^2$
$R_{pop}^2 := \frac{E(\beta'X)^2}{EY^2} =$	$R_{sample}^2 := \frac{\mathbb{E}_n(\hat{\beta}'X_i)^2}{\mathbb{E}_n Y_i^2} =$
$1 - \frac{E\varepsilon^2}{EY^2} \in [0, 1]$	$1 - \frac{\mathbb{E}_n \hat{\varepsilon}_i^2}{\mathbb{E}_n Y_i^2} \in [0, 1]$

By law of large numbers when p/n is small and n is large:

$$\begin{aligned} \mathbb{E}_n Y_i^2 &\approx EY^2, \quad \mathbb{E}_n(\hat{\beta}'X_i)^2 \approx E(\beta'X)^2, \quad \mathbb{E}_n \hat{\varepsilon}_i^2 \approx E\varepsilon^2 \\ R_{sample}^2 &\approx R_{pop}^2 \quad \text{and} \quad MSE_{sample} \approx MSE_{pop} \end{aligned} \tag{3}$$

Overfitting: What happens when p/n is not small

When p/n is not small, the discrepancy between the in-sample and out-of-sample measures of fit can be substantial. Let's check the next example :

-

$$\begin{aligned} X &\sim N(0, I_p) \text{ and } Y \sim N(0, 1), \beta'X = 0, R_{pop}^2 = 0 \\ \text{if } p &= n, \text{ then } R_{sample}^2 \text{ is } 1 \gg 0 \\ \text{if } p &= \frac{n}{2}, \text{ then } R_{sample}^2 \text{ is about } 0.5 \gg 0 \\ \text{if } p &= \frac{n}{20}, \text{ then } R_{sample}^2 \text{ is about } 0.05 \end{aligned} \tag{4}$$

Better measures of out-of-sample predictive ability are the “adjusted” R^2 and MSE .

$$MSE_{adjusted} = \frac{n}{n-p} \mathbb{E}_n \hat{\epsilon}_i^2, \quad R_{adjusted}^2 := 1 - \frac{n}{n-p} \frac{\mathbb{E}_n \hat{\epsilon}_i^2}{\mathbb{E}_n Y_i^2} \tag{5}$$

Measuring Predictive Ability by Sample Splitting

To measure out-of-sample performance: **Data splitting**. The idea can be summarized in two parts:

1. Use a random part of data, called the **training sample**, for estimating/training the prediction rule.
2. Use the other part, called the **testing sample**, to evaluate the quality of the prediction rule, recording out-of-sample mean squared error and R^2 .

Generic Evaluation of Prediction Rules by Sample-Splitting

1. Randomly partition the data into training and testing samples. Suppose we use n observations for training and m for testing/validation.
2. Use the training sample to compute a prediction rule $\hat{f}(X)$, for example, $\hat{f}(X) = \beta'X$.
3. Let V denote the indexes of the observations in the test sample. Then the out-of-sample/test mean squared error is

$$MES_{test} = \frac{1}{m} \sum_{k \in V} (Y_k - \hat{f}(X_k))^2 \quad (6)$$

and the out-of-sample/test R^2 is

$$R_{test}^2 = 1 - \frac{MSE_{test}}{\frac{1}{m} \sum_{k \in V} Y_k^2} \quad (7)$$

Understanding β_1 via "Partialling-Out"

1. Lets imagine we have the next equation

$$Y = \underbrace{[\beta_1 D + \beta_2' W]}_{\text{Predicted value}} + \underbrace{\varepsilon}_{\text{error}} \quad (8)$$

How does the predicted value of Y change if D increases by a unit while W remains unchanged?

2. **Partialling-out operation:** procedure that takes a random variable V and creates a "residual" \tilde{V} by subtracting the part of V that is linearly predicted by:

$$\tilde{V} = V - \gamma'_{VW} W, \quad \gamma_{VW} = \arg \min_{\gamma} E(V - \gamma' W)^2 \quad (9)$$

3. We can show that

$$Y = V + U \implies \tilde{Y} = \tilde{V} + \tilde{U} \quad (10)$$

Understanding β_1 via "Partialling-Out"

1. Partialling-out to both sides of our regression equation $Y = \underbrace{[\beta_1 D + \beta_2' W]}_{\text{Predicted value}} + \underbrace{\varepsilon}_{\text{error}}$, we get:

$$\tilde{Y} = \beta_1 \tilde{D} + \beta_2' \tilde{W} + \tilde{\varepsilon} \quad (11)$$

which simplifies to :

$$\tilde{Y} = \beta_1 \tilde{D} + \varepsilon, E[\varepsilon \tilde{D}] = 0 \quad (12)$$

- $\beta_2' \tilde{W} = 0$
- $\tilde{\varepsilon} = \varepsilon$
- $E[\varepsilon \tilde{D}] = 0$; since \tilde{D} is a linear function of $X = (D, W)$

2. What we found in (16) are the Normal Equations for the population regression of \tilde{Y} on \tilde{D} .

Understanding β_1 via "Partialling-Out"

Theorem (Frisch-Waugh-Lovell, FWL)

The population linear regression coefficient β_1 can be recovered from the population linear regression of \tilde{Y} on \tilde{D} .

$$\beta_1 = \arg \min_{b_1} E(\tilde{Y} - b_1 \tilde{D})^2 = (E\tilde{D}^2)^{-1} E\tilde{D}\tilde{Y}, \quad E\tilde{D}^2 > 0 \quad (13)$$

Theorem (Frisch-Waugh-Lovell, FWL)

Follow the next algorithm:

- Regress Y on W , obtain residuals ε_1
- Regress D on W , obtain residuals ε_2
- Regress ε_1 on ε_2 , obtain OLS estimates β_1

Understanding β_1 via "Partialling-Out"

Theorem (Frisch-Waugh-Lovell, FWL)

*In other words, β_1 can be interpreted as a (univariate) linear regression coefficient in the linear regression of **residualized Y on residualized D** , where the residuals are defined by **partialling-out the linear effect of W from Y and D** .*

Understanding $\hat{\beta}_1$ via "Partialling-Out"

Theorem (Frisch-Waugh-Lovell, FWL-IN SAMPLE)

When we work with the sample, we mimic the partialling out in the population.

$$\hat{\beta}_1 = \arg \min_{b_1} \mathbb{E}(\check{Y}_i - \beta_1 \check{D}_i)^2 = (\mathbb{E} \check{D}_i^2)^{-1} \mathbb{E} \check{D}_i \check{Y}_i \quad (14)$$

where \check{Y}_i denote the residual left after predicting V_i with controls W_i in the sample:

$$\check{Y}_i = V_i - \hat{\gamma}'_{VW} W_i, \quad \hat{\gamma}_{VW} = \arg \min_{\gamma} \mathbb{E}_n(V_i - \gamma' W_i)^2 \quad (15)$$

Understanding β_1 via "Partialling-Out"

1. Frisch-Waugh-Lovell-Python
2. Frisch-Waugh-Lovell-Julia

Regression in a High-Dimensional Setting / LASSO

Lasso constructs $\hat{\beta}$ as the solution of the following penalized least squares problem:

$$\min_{b \in \mathbb{R}^p} \sum_i (Y_i - b'X_i)^2 + \lambda \cdot \sum_{j=1}^p |b_j| \quad (16)$$


1. The first term is n times the sample mean square error
2. The second term is a penalty term, which penalizes the size of coefficients b_j by their absolute values times the penalty level λ
A crucial point is the choice of the penalization parameter λ .
3. A theoretically valid choice is (Belloni Chernozhukov, 2013)

$$\lambda = 2 \cdot c \hat{\sigma} \sqrt{2n \log(2p/\gamma)}, \quad \hat{\sigma} \approx \sigma = \sqrt{E \epsilon^2} \quad (17)$$

4. Another good way to pick penalty level is by cross-validation (Chetverikov et al, 2020)

Contours of the error and constraint functions for the lasso

Figure 1: Lasso optimization with two coefficients.



Intuition : The j -th component $\hat{\beta}_j$ of the lasso estimator $\hat{\beta}$ is set to zero if the marginal predictive benefit of changing $\hat{\beta}_j$ away from zero is smaller than the marginal increase in penalty:

$$\hat{\beta}_j = 0 \text{ if } \left| \partial_{b_j} \sum_i (Y_i - \hat{\beta}' X_i)^2 \right| < \lambda \quad (18)$$

OLS post Lasso

We can then **use the Lasso-selected set of regressors** to refit the model by least squares. This method is called the “least squares post Lasso” or simply **post-Lasso**

$$\tilde{\beta} \in \arg \min_{\beta \in \mathbb{R}^p} \sum_i (Y_i - X_i' \beta)^2 : \beta_j = 0 \text{ if } \hat{\beta}_j = 0 \text{ for each } j \quad (19)$$

Under approximate sparsity Lasso and Post-Lasso will approximate the best linear predictor well. This means that they won't overfit the data, and we can use the sample and adjusted R^2 and MSE to assess out-of-sample predictive performance. Of course, it is always a good idea to verify the out-of-sample predictive performance by using sample splitting

How to select λ

Big lambdas tend to result in a lot of shrinkage and sparsity, as $\lambda \rightarrow 0$ our solution approaches the OLS solution

Two ways to select λ

- Select model with lowest AIC/BIC/other plug-in criterion. This uses no out-of-sample information for selection but is fast.
- Cross-validate by testing on our hold-out test sample. **Variants of cross-validation are most commonly used.**

k-fold cross-validation

The next algorithm is taken from Ivan Rudik' Lectures Note in Dynamic Optimization(Cornell, Fall 2021)

Figure 2: k-fold cross-validation

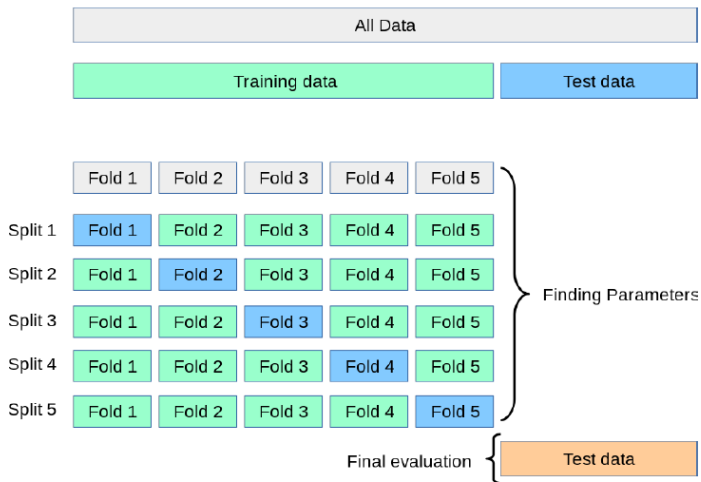
k-fold cross-validation

In k-fold cross-validation we do the following:

- Create a grid of λ s
- For each λ :
 - Split data into k mutually-exclusive folds of about equal size, usually choose $k = 5, 10$
 - For $j = 1, \dots, k$
 - fit the model using all folds but fold j
 - Predict out-of-sample on fold j
 - Compute average mean squared prediction error across the k folds:
$$\bar{Q}(\lambda) = \frac{1}{k} \sum_{j=1}^k \sum_{i \in \text{fold } j} (y_i - (\alpha_0 + x'_i \beta))^2 + \lambda \|\beta\|_1$$
- Choose $\hat{\lambda}_{min} = \operatorname{argmin}_{\lambda} \bar{Q}(\lambda)$ or to avoid modest overfitting choose the largest λ such that $\bar{Q}(\lambda) \leq \hat{\lambda}_{min} + \sigma_{\hat{\lambda}_{min}}$ (1 standard deviation rule)

k-fold cross-validation

Figure 3: k-fold cross-validation



Practical Value of CEF

1. Randomly split your sample with N observations into a training set with N_{tr} and a test set with N_{te} observations.
2. Estimate different CEF functions in the training sample $\hat{m}(X)$.
3. Calculate and compare their out-of-sample Mean Squared Error (MSE) in the test sample:

$$MSE_{te} = \frac{1}{N_{te}} \sum_{i=1}^{N_{te}} (Y_i - \hat{m}(X_i))^2$$

4. Pick the function with the lowest MSE.